

Chapter 4

Network Layer:

Data Plane

A note on the use of these PowerPoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR


All material copyright 1996-2020
J.F Kurose and K.W. Ross, All Rights Reserved



*Computer Networking: A
Top-Down Approach*
8th edition
Jim Kurose, Keith Ross
Pearson, 2020

Network-layer services and protocols

- transport segment from sending to receiving host
 - **sender:** encapsulates segments into datagrams, passes to link layer
 - **receiver:** decapsulates segments from datagrams, delivers them to transport layer
- network layer protocols in *every Internet device*: hosts, routers
- **routers:**
 - examines header fields in all datagrams passing through it
 - moves datagrams from input ports to output ports to transfer datagrams along end-end path



Two key network-layer functions

network-layer functions:

- *forwarding*: move packets from a router's input link to appropriate router output link
- *routing*: determine route taken by packets from source to destination
 - *routing algorithms and protocols*

analogy: taking a trip

- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination



forwarding




routing

Network layer: data plane, control plane

Data plane:

- *local*, per-router function
- determines how datagram arriving on router input port is forwarded to router output port
 - implements *forwarding*




Control plane:

- *network-wide* logic
- determines how datagram is routed among routers along end-end path from source host to destination host
 - implements *routing*
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers

Network layer: data plane, control plane

Data plane:

- *local*, per-router function
- determines how datagram arriving on router input port is forwarded to router output port
 - implements *forwarding*




Control plane:

- *network-wide* logic
- determines how datagram is routed among routers along end-end path from source host to destination host
 - implements *routing*
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers

out of scope


Traditional routing: per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane




Software-Defined Networking (SDN) control plane

Remote controller computes, installs forwarding tables in routers



Software-Defined Networking (SDN) control plane

Remote controller computes, installs forwarding tables in routers



Network service model

Q: What *service model* for “channel” transporting datagrams from sender to receiver?

example services for
individual datagrams:

- guaranteed delivery
- guaranteed delivery with less than 40 msec delay
- ...

example services for a
flow of datagrams:

- in-order datagram delivery
- guaranteed minimum bandwidth to flow
- ...

Network-layer service model

Network Architecture	Service Model	Quality Guarantees?		
		Loss	Order	Bandwidth
Internet	<i>best effort</i>	no	no	no


Internet “best effort” service model

No guarantees on:


- i. successful datagram delivery to destination
- ii. order of delivery
- iii. bandwidth available to end-end flow

Router architecture overview

high-level view of generic router architecture:




Input port functions



- using header field values, lookup output port using forwarding table in input port memory (*"match plus action"*)
- **destination-based forwarding:** forward based only on destination network-layer (i.e., IP) address (traditional)
 - We will see how it works in a few slides...

Input port functions




decentralized switching/forwarding:

- using header field values, lookup output port using forwarding table in input port memory ("*match plus action*")
- goal: complete input port processing in a very short time
- **input port queuing:** if datagrams arrive faster than *switching rate* of switching fabric


Switching fabric

- transfer packet from input port to appropriate output port
- **switching rate:** rate at which packets can be transferred from input ports to output ports
 - often measured as multiple of input/output port line rate (i.e., rate supported by the port)
 - N inputs: switching rate at least N times line rate R desirable → negligible input port queuing is ensured




Input port queuing

- if switching rate slower than input ports combined → significant queueing may occur at input queues
 - queueing delay and loss due to input buffer overflow!
- Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward




output port contention: only one red datagram can be transferred; lower red packet is *blocked*




one packet time later: green packet has experienced *HOL blocking*

Output port queuing




- Buffering occurs when arrival rate from the switching fabric exceeds output line speed
 - *Queueing delay and loss due to output port buffer overflow!*

Output port queuing




- Buffering occurs when arrival rate from the switching fabric exceeds output line speed
 - *Queueing delay and loss due to output port buffer overflow!*
- *Drop policy:* which datagrams to drop if buffers are almost full? → Datagrams are discarded
- *Scheduling discipline* chooses among queued datagrams for transmission → Priority scheduling – who gets best performance


IP Datagram format




IP Datagram format




IP Datagram format




IP Datagram format




IP Datagram format




IP Datagram format




IP Datagram format




IP Datagram format




IP Datagram format



IP Datagram format




IP Datagram format



IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
 - router's typically have multiple interfaces
 - *port* and *interface* are synonyms in our context
 - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)



dotted-decimal IP address notation:

223.1.1.1 =


11011111	00000001	00000001	00000001
223	1	1	1

IP addressing: introduction

Q: how are interfaces actually connected?

A: we'll partially learn about that in chapter 6

A: wired Ethernet interfaces connected by Ethernet switches



For now: don't need to worry about how one interface is connected to another (with no router in between)

A: wireless WiFi interfaces connected by WiFi base station (not investigated)


Subnets

- *What's a subnet?*

- device interfaces that can reach each other **without passing through a router**

- IP addresses have a structure:

- **subnet part:** devices in same subnet have common high order bits
- **host part:** remaining low order bits



How many subnets in this case?


Subnets

- *What's a subnet?*

- device interfaces that can reach each other **without passing through a router**

- IP addresses have a structure:

- **subnet part:** devices in same subnet have common high order bits
- **host part:** remaining low order bits




How many subnets in this case?

Network consisting of 3 subnets

Subnets

Recipe for defining subnets:

- detach each interface from its host or router, creating “islands” of isolated networks
- each isolated network is called a *subnet*
- what do the **/24** subnet addresses mean?
 - They are the *subnet addresses* (all 0s for the host part)



IP addressing: Classless (CIDR)

CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
 - address format: $a.b.c.d/x$, where x is # bits in subnet portion of address



Alternative representation of /23: **Subnet Mask**

11111111 111111111 111111110 00000000
255.255.254.0

IP addressing: Classful

First byte


Class A
(0-127)



Class B
(128-191)




Class C
(192-223)



- No need for subnet mask
 - Subnet and host portions can be identified by looking at the first bits of subnet portion
- **Not used anymore**
 - Addressing is too rigid!

Special IP addresses

- **Subnet address:** host part with all zeros
 - E.g. 193.17.31.0/24
- **Direct broadcast address:** host part with all ones
 - E.g. 193.17.31.255
- **Limited broadcast address:** all ones (255.255.255.255)
 - Broadcast in the same subnet
 - The message cannot overcome routers



Destination-based forwarding

Destination-based forwarding

Example of forwarding table

	Destination IP Address Range				Link interface
Subnet mask: \21	11001000	00010111	00010***	*****	0
Subnet mask: \24	11001000	00010111	00011000	*****	1
Subnet mask: \21	11001000	00010111	00011***	*****	2
	otherwise				3

Destination-based forwarding

Example of forwarding table

	Destination IP Address Range				Link interface
Subnet mask: \21	11001000	00010111	00010***	*****	0
Subnet mask: \24	11001000	00010111	00011000	*****	1
Subnet mask: \21	11001000	00010111	00011***	*****	2
	otherwise				3

examples:

11001000 00010111 00010110 10100001 which interface?
11001000 00010111 00011000 10101010 which interface?

Destination-based forwarding

Destination IP Address Range	Link interface
11001000 00010111 00010*****	0
11001000 00010111 00011000 *****	1
11001000 1 00011*** *****	2
otherwise	3

examples:

11001000 00010111 00010110 10100001 which interface?

11001000 00010111 00011000 10101010 which interface?

Destination-based forwarding

Destination IP Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

match!

11001000 00010111 00010110 10100001
11001000 00010111 00011000 10101010

which interface?

which interface?

examples:

Destination-based forwarding

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

chosen!

match!

examples:

Destination IP Address Range					Link interface
11001000	00010111	00010***	*****	0	
11001000	00010111	00011000	*****	1	
11001000	00010111	00011***	*****	2	
otherwise				3	

11001000	00010111	00010110	10100001	which interface?
11001000	00010111	00011000	10101010	which interface?

IP addresses: how to get one?

That's actually **two** questions:

- 1. Q:** How does a *host* get IP address within its network (host part of address)?
- 2. Q:** How does a *network* get IP address for itself (network part of address)

How does a *host* get IP address?

- Statically specified in config file of the OS (e.g., `/etc/rc.config` in UNIX)
- **DHCP: Dynamic Host Configuration Protocol** → dynamically get address from a server (called *DHCP server*)
 - “plug-and-play”

DHCP: Dynamic Host Configuration Protocol


goal: host *dynamically* obtains IP address from network server when it “joins” the network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/on)

DHCP overview:

- host broadcasts **DHCP discover** msg [*optional*]
- DHCP server responds with **DHCP offer** msg [*optional*]
- host requests IP address: **DHCP request** msg
- DHCP server sends address: **DHCP ack** msg

DHCP client-server scenario




Typically, DHCP server will be co-located in router, serving subnets to which router is attached

arriving **DHCP client** needs address in this network

DHCP client-server scenario


DHCP server: 223.1.2.5



Arriving client



DHCP client-server scenario



DHCP client-server scenario


DHCP server: 223.1.2.5



DHCP discover


```
src IP, port: 0.0.0.0, 68  
dest IP, port: 255.255.255.255,67  
assignedIPaddr: 0.0.0.0  
transaction ID: 654
```

Arriving client



DHCP client-server scenario

DHCP server: 223.1.2.5



DHCP client-server scenario

DHCP server: 223.1.2.5



DHCP discover


```
src IP, port: 0.0.0.0, 68  
dest IP, port: 255.255.255.255,67  
assignedIPaddr: 0.0.0.0  
transaction ID: 654
```

Arriving client



DHCP offer

```
src IP, port: 223.1.2.5, 67  
dest IP, port: 255.255.255.255, 68  
assignedIPaddr: 223.1.2.4  
transaction ID: 654  
lifetime: 3600 secs
```



DHCP client-server scenario

DHCP server: 223.1.2.5



DHCP discover

```
src IP, port: 0.0.0.0, 68  
dest IP, port: 255.255.255.255,67  
assignedIPaddr: 0.0.0.0  
transaction ID: 654
```


Arriving client



DHCP offer


```
src IP, port: 223.1.2.5, 67  
dest IP, port: 255.255.255.255, 68  
assignedIPaddr: 223.1.2.4  
transaction ID: 654  
lifetime: 3600 secs
```

The two steps above can be skipped “if a client remembers and wishes to reuse a previously allocated network address”
[RFC 2131]




DHCP client-server scenario

DHCP server: 223.1.2.5




DHCP client-server scenario

DHCP server: 223.1.2.5




DHCP client-server scenario

DHCP server: 223.1.2.5



DHCP client-server scenario

DHCP server: 223.1.2.5



DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:

- name and IP address of the **local DNS server**
- **subnet mask** (indicating network versus host portion of address)
- **default gateway** (IP address of first-hop router)

IP addresses: how to get one?

Q: how does *network* get subnet part of IP address?

A: gets allocated portion of its provider ISP's address space

ISP's block 11001000 00010111 00010000 00000000 200.23.16.0/20

Let's assume that there are 8 organizations that need to get a subnet address → ISP can then split its address space in 8 blocks:

IP addresses: how to get one?

Q: how does *network* get subnet part of IP address?

A: gets allocated portion of its provider ISP's address space

ISP's block 11001000 00010111 00010000 00000000 200.23.16.0/20

Let's assume that there are 8 organizations that need to get a subnet address → ISP can then split its address space in 8 blocks:

Organization 0 11001000 00010111 00010000 00000000 200.23.16.0/23

Organization 1 11001000 00010111 00010010 00000000 200.23.18.0/23

Organization 2 11001000 00010111 00010100 00000000 200.23.20.0/23

...

.....

.....

....

Organization 7 11001000 00010111 00011110 00000000 200.23.30.0/23

IP addresses: how to get one?

Q: how does *network* get subnet part of IP address?

A: gets allocated portion of its provider ISP's address space

ISP's block 11001000 00010111 00010000 00000000 200.23.16.0/20

Let's assume that there are 8 organizations that need to get a subnet address → ISP can then split its address space in 8 blocks:

Organization 0 11001000 00010111 00010000 00000000 200.23.16.0/23

Organization 1 11001000 00010111 00010010 00000000 200.23.18.0/23

Organization 2 11001000 00010111 00010100 00000000 200.23.20.0/23

...

.....

.....


....

Organization 7 11001000 00010111 00011110 00000000 200.23.30.0/23

This operation is called *subnetting* or *hierarchical addressing*


Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information (*route aggregation*):




Hierarchical addressing: more specific routes

- Organization 1 moves from ISP 1 to ISP 2
- ISP 2 now advertises a more specific route to Organization 1




Hierarchical addressing: more specific routes

- Organization 1 moves from ISP 1 to ISP 2
- ISP 2 now advertises a more specific route to Organization 1




Hierarchical addressing: more specific routes

- Organization 1 moves from ISP 1 to ISP 2
- ISP 2 now advertises a more specific route to Organization 1




Hierarchical addressing: more specific routes

- Organization 1 moves from ISP 1 to ISP 2
- ISP 2 now advertises a more specific route to Organization 1




Hierarchical addressing: more specific routes

- Organization 1 moves from ISP 1 to ISP 2
- ISP 2 now advertises a more specific route to Organization 1



Hierarchical addressing: more specific routes

- Organization 1 moves from ISP 1 to ISP 2
- ISP 2 now advertises a more specific route to Organization 1



IP addressing: last words...

Q: how does an ISP get block of addresses?

A: ICANN: Internet Corporation for Assigned Names and Numbers
<http://www.icann.org/>


Q: are there enough 32-bit IP addresses?

- ICANN allocated last chunk of IPv4 addresses in 2011
- IPv6 has 128-bit address space
- NAT (next) helps IPv4 address space exhaustion

“Who the hell knew how much address space we needed?” - *Vint Cerf* (reflecting on decision to make IPv4 address 32 bits long)

NAT: network address translation

NAT: all devices in local network share just **one** IPv4 address as far as outside world is concerned



all datagrams *leaving* local network have *same* source NAT IP address: 138.76.29.7, but *different* source port numbers

datagrams with source or destination in this network have 10.0.0.0/8 address for source, destination (as usual)

NAT: network address translation

- all devices in local network have 32-bit addresses in a “private” IP address space (10.0.0.0/8, 172.16.0.0/12, 192.168.0.0/16 prefixes) that can only be used in local network
- advantages:
 - just **one** “public” IP address needed from provider ISP for ***all*** devices
 - can change addresses of host in local network without notifying outside world
 - can change ISP without changing addresses of devices in local network
 - security: devices inside local net not directly addressable, visible by outside world


NAT: network address translation

implementation: NAT router must (transparently):

- outgoing datagrams: replace (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
 - remote clients/servers will respond using (NAT IP address, new port #) as destination address
- remember (in **NAT translation table**) every (source IP address, port #) to (NAT IP address, new port #) translation pair
- incoming datagrams: replace (NAT IP address, new port #) in destination fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: network address translation


NAT translation table	
outside addr	LAN side addr



NAT: network address translation


NAT translation table	
outside addr	LAN side addr

1: host 10.0.0.1 sends datagram to 128.119.40.186, 80




NAT: network address translation


2: NAT router changes datagram source address from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table



NAT: network address translation



NAT: network address translation



NAT: network address translation

- NAT has been controversial:
 - routers “should” only process up to layer 3
 - violates end-to-end argument (port # manipulation by network-layer device)
 - address “shortage” should be solved by IPv6
- but NAT is here to stay:
 - extensively used in home and institutional nets, 4G/5G cellular nets